

Primary sources & further reading

Part 12 —  AI Governance · 12-ai-governance/papers/PAPERS.md · 12-ai-governance/resources/RESOURCES.md



Papers & Primary Sources — AI Governance

Unlike most modules in this track, the primary sources here are **frameworks and standards documents**, not papers with a novel method. Read the frameworks first — everything else in this module operationalizes them. A handful of research papers underpin the concrete artifacts (model cards, datasheets) and the risk taxonomy this module's risk register is built from; those are worth reading in full too, they're short.

Read these two first

#	Source	Why it matters	Link
1	NIST AI Risk Management Framework 1.0 (NIST AI 100-1, Jan 2023)	The four-function structure (Govern/Map/Measure/Manage) this module's notes/01 §2 is built on. Free, voluntary, and the closest thing to a lingua franca for AI risk in the US.	PDF · NIST landing page
2	Model Cards for Model Reporting — Mitchell, Wu, Zaldivar, Barnes, Vasserman, Hutchinson, Spitzer, Raji, Gebru (FAT* 2019)	The paper that invented the model card as a governance artifact. <code>code/model_card_template.md</code> is a direct descendant of this structure (intended use, factors, metrics, evaluation data, training data, ethical considerations, caveats).	arXiv:1810.03993

Frameworks and standards (the core reading)

Source	Publisher / Year	One-line takeaway	Link
AI RMF Playbook	NIST, 2023 (updated periodically)	Suggested actions per AI RMF subcategory — the practical companion to the framework itself; adapt, don't copy verbatim.	AIRC Play book
NIST AI 600-1 — Generative AI Profile	NIST, July 2024	Maps the four AI RMF functions to 12 risks specific to or amplified by generative AI (confabulation, data privacy, information integrity, value-chain integration, and others). Read this if your use case involves an LLM, which in this track is almost always.	PDF
ISO/IEC 42001:2023 — AI management systems	ISO/IEC, Dec 2023	The certifiable AI management system standard behind notes/01 §3 — Harmonized Structure, clauses 4-10, Annex A's 38 controls across 9 groups, Statement of Applicability. The standard text itself is paywalled; the overview page is authoritative and free.	ISO over view
SR 11-7 — Supervisory Guidance on Model Risk Management	Federal Reserve / OCC, 2011	The banking-sector origin of the model risk management (MRM) discipline this module borrows in notes/02 §2 — independent validation, separate from model development, as a structural control. Still the reference text bank examiners cite.	PDF
OECD AI Principles	OECD, 2019, updated May 2024	The first intergovernmental AI standard (47 adherents as of 2024); five principles — human-centered values, transparency, robustness/safety, accountability, inclusive growth — that most national frameworks (including the EU AI Act's recitals) trace back to.	OEC D.AI
Singapore Model AI Governance Framework (+ Generative AI and Agentic AI extensions)	IMDA / PDPC, 2019-2026	A genuinely practical, non-binding, proportionality-first governance framework — worth reading specifically because it's less legalistic than the EU AI Act and closer in spirit to this module's internal-process focus. The 2026 Agentic AI extension is a useful bridge to Module 07 [07-agentic-ai].	PDP C

Foundational documentation-transparency papers

Paper	Year	One-line takeaway	Link
Datasheets for Datasets — Gebru, Morgenstern, Vecchione, Vaughan, Wallach, Daumé III, Crawford	2018 (rev. 2021)	The dataset-side sibling of the model card: motivation, composition, collection process, recommended and discouraged uses. <code>code/model_card_template.md</code> §4's "training/grounding data summary" is a compressed datasheet.	arXiv: 1803. 09010
"Model Cards for Model Reporting" in 2024: Reclassifying Ethical Considerations — various	2024	A five-years-later retrospective reframing the original model card's "ethical considerations" section in terms of trustworthiness and risk management — useful for seeing how the artifact has aged.	arXiv: 2403. 15394
GPT-4 System Card — OpenAI	2023	The most-cited real-world example of a system card (vs. model card) — explicitly analyzes the deployed system including non-model mitigations (usage policies, access controls), which is exactly the distinction notes/02 §3 draws for the RAG example.	PDF

Risk taxonomy and incident evidence

Paper / resource	Year	One-line takeaway	Link
Ethical and Social Risks of Harm from Language Models — Weidinger et al. (DeepMind)	2021 (FAccT '22 companion)	21 risks in 6 categories (discrimination, information hazards, misinformation, malicious use, human-computer interaction harms, environmental/socioeconomic harms) — the taxonomy most "what could go wrong" governance intake forms are quietly built from.	arXiv:2112.04359
AI Incident Database (AIID)	Partnership on AI / Responsible AI Collaborative, ongoing since 2020	1,200+ real, sourced AI incidents. The single best antidote to "that would never happen here" in a governance committee meeting — pull three incidents in your sector before your next risk review.	incidentdatabase.ai

How to read a governance framework document (30 minutes)

Frameworks aren't papers — read them differently:

1. **Pass 1 (10 min):** scope and unit of analysis. Is this voluntary guidance, a certifiable standard, or binding law? Who is it written for (a whole org, a single system, a specific sector)?
2. **Pass 2 (15 min):** the structure itself — the functions/clauses/categories. Map each one to something your organization already does or is missing; that gap list is your adoption backlog.
3. **Pass 3 (5 min):** what evidence would prove you're actually doing this, not just intending to. If the framework doesn't suggest an artifact (a document, a log, a sign-off), it's a value statement, not a control — note it as aspirational and move on.

Keep one paragraph per framework: what it's for, what it isn't for, and which of your existing artifacts already partially satisfies it. In six months that paragraph is worth more than the PDF.

Resources — AI Governance

Frameworks and standards bodies (the primary references)

Framework	Publisher	What it's for	Link
AI Risk Management Framework (AI RMF 1.0)	NIST	Voluntary four-function structure (Govern/Map/Measure/Manage) — the fastest free starting point for any org	https://www.nist.gov/itl/ai-risk-management-framework
AI RMF Playbook	NIST	Suggested actions per subcategory	https://airc.nist.gov/airmf-resources/playbook
NIST AI 600-1 — Generative AI Profile	NIST	12 generative-AI-specific risk categories mapped onto the AI RMF	https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.600-1.pdf
ISO/IEC 42001:2023	ISO/IEC	The certifiable AI management system (AIMS) standard	https://www.iso.org/standard/42001
SR 11-7 — Model Risk Management	Federal Reserve / OCC	Origin of the MRM discipline this module borrows for independent model validation	https://www.federalreserve.gov/boarddocs/srletters/2011/sr1107.pdf
OECD AI Principles	OECD	The intergovernmental baseline most national frameworks cite	https://oecd.ai/en/ai-principles
Singapore Model AI Governance Framework (+ Gen AI / Agentic AI extensions)	IMDA / PDPC	A proportionality-first, non-binding framework — good contrast to the EU AI Act's binding tiers	https://www.pdpc.gov.sg/help-and-resources/2020/01/model-ai-governance-framework
AI Verify	IMDA (Singapore)	An actual open-source testing toolkit/framework for validating AI system claims against governance principles, not just a document	https://aiverifyfoundation.sg/

Organizations and communities

Org	What they do	Link
Partnership on AI	Multi-stakeholder nonprofit; runs the AI Incident Database and the ABOUT ML documentation-standards project	https://partnershiponai.org/
AI Incident Database (AIID)	1,200+ sourced, searchable real-world AI failure reports — the best "this actually happened" evidence for a risk committee	https://incidentdatabase.ai/
NIST AI Resource Center (AIRC)	Hub for all NIST AI RMF materials, crosswalks (e.g. NIST 600-1 vs. IMDA AI Verify), and profiles	https://airc.nist.gov/
OECD.AI Policy Observatory	Tracks national AI policy and governance initiatives across 47+ countries	https://oecd.ai/
ForHumanity	Independent nonprofit building auditable, certifiable AI governance criteria (used in some EU AI Act conformity assessments)	https://forhumanity.center/

Cloud-provider responsible AI programs (useful for vendor-risk review, notes/02 §5)

Provider	Program	Link
Google Cloud	Responsible AI practices, Vertex AI model cards, Frontier Safety Framework	https://cloud.google.com/responsible-ai · https://ai.google/responsibility/
Microsoft	Responsible AI Standard + per-service Transparency Notes (Azure AI Foundry)	https://www.microsoft.com/en-us/ai/responsible-ai · https://learn.microsoft.com/en-us/azure/ai-foundry/responsible-ai/openai/transparency-note
AWS	Responsible AI policy, AWS Well-Architected Responsible AI Lens, Skill Builder courses	https://aws.amazon.com/ai/responsible-ai/policy/ · https://aws.amazon.com/machine-learning/responsible-ai/

When you're assessing a vendor per notes/02 §5, these pages are exactly the kind of evidence to ask for — "can you point me to your model/system card and your responsible-AI documentation for the specific model version we'd be calling" is a fair, answerable question precisely because these programs exist.

Courses (free or free-to-audit)

Course	Who it's for	Link
AWS Skill Builder — Introduction to Responsible AI / Security, Compliance & Governance for AI Solutions	Practitioners who want a vendor's-eye view of governance controls	https://skillbuilder.aws/
Google — Responsible AI: Applying AI Principles with Google Cloud	Product/eng teams shipping on Vertex AI	https://www.skills.google/course_templates/388
Google for Developers — Introduction to Responsible AI	Quick, free, conceptual grounding	https://developers.google.com/machine-learning/guides/intro-responsible-ai
Coursera — AI Governance / Responsible AI courses (search current catalog; offerings rotate)	Structured, credentialed option if your org wants a certificate	https://www.coursera.org/

Tools worth knowing about

Tool	What it does	Link
AI Verify	Open-source testing framework/toolkit to validate AI system claims (fairness, robustness, transparency) against governance principles	https://aiverifyfoundation.sg/
<code>code/ai_risk_register_template.py</code> (this module)	The risk-tiering mechanic implemented as a runnable, adjustable tool — start here before reaching for anything heavier	<code>../code/ai_risk_register_template.py</code>
Model card generation tooling (Hugging Face Hub model cards)	A widely-used, real implementation of the Mitchell et al. model card structure, worth studying as a working example	https://huggingface.co/docs/hub/en/model-cards

Cross-references inside this track

- Module 13 — AI Security [13-ai-security] — the technical controls a governance policy points at (prompt injection defenses, red teaming).
- Module 14 — AI Compliance [14-ai-compliance] — the external, binding-law side of the same risk-tiering logic (EU AI Act, sector regulation).
- Module 15 — AI Evals [15-ai-evals] — where a model/system card's evaluation section actually comes from.
- Module 16 — LangGraph track [16-ai-tech-stack/tracks/langgraph] — the technical mechanism (`interrupt()` / `Command(resume=...)`) that implements a human-oversight policy in code.

Communities

- Partnership on AI mailing list / workstreams — <https://partnershiponai.org/>
- r/artificial and r/MachineLearning (governance and policy threads surface regularly) — <https://www.reddit.com/r/artificial/>
- OECD.AI network of experts — <https://oecd.ai/en/network-of-experts>