

Primary sources & further reading

Part 14 —  AI Compliance · 14-ai-compliance/papers/PAPERS.md · 14-ai-compliance/resources/RESOURCES.md



Papers & Primary Sources — AI Compliance

Read the primary legal texts first, commentary second. Law-firm blog posts and LinkedIn summaries lag the actual text and compound each other's errors — a pattern this module warns about repeatedly. Every link below was checked at the time this file was written (September 2026); re-verify anything date-sensitive per notes/03's "how to stay current" [14-ai-compliance/notes/03-building-a-compliance-program.md].

Read these two first

#	Source	Why it matters	Link
1	Regulation (EU) 2024/1689 (the EU AI Act) — official consolidated text	The actual binding law. Everything in notes/01 is a summary of this; cite this, never the summary, for anything load-bearing.	EUR-Lex
2	NIST AI RMF 1.0 (NIST AI 100-1, Jan 2023)	The most-referenced US framework — voluntary, but cited by the FTC, CFPB, FDA, SEC, and EEOC. The Govern/Map/Measure/Manage structure this module's US material assumes.	PDF

The EU AI Act — text, guidance, and the Digital Omnibus

Source	What it is	Link
Regulation (EU) 2024/1689 — full text, all languages	The primary source	EUR-Lex
EU AI Act Service Desk	European Commission's official Q&A and guidance portal, updated as amendments land	ai-act-service-desk.ec.europa.eu
Implementation Timeline tracker	Independently maintained but closely tracks official dates; the single best place to check "what's in force right now"	artificialintelligenceact.eu/implementation-timeline
European Commission — AI Act policy page	The Commission's own framing of the regulatory approach	digital-strategy.ec.europa.eu
GPAI Code of Practice	The Commission's voluntary code for general-purpose-AI providers (transparency, copyright, safety/security)	digital-strategy.ec.europa.eu
Digital Omnibus on AI (2026 amendment)	The 2026 package that deferred several high-risk deadlines — read this to understand <i>why</i> notes/01's timeline table has "deferred" rows, not just <i>that</i> it does	Search EUR-Lex for the Digital Omnibus's official regulation number at time of reading; the number was still being finalized in secondary sources as of this module's last check — confirm via the Service Desk link above rather than trusting a specific citation here

NIST AI Risk Management Framework and profiles

Source	What it is	Link
NIST AI RMF 1.0 (NIST AI 100-1)	The core framework	PDF
NIST AI RMF: Generative AI Profile (NIST AI 600-1, July 2024)	Adapts Govern/Map/Measure/Manage to generative AI and agentic systems; 12 named risk categories	PDF
NIST AI RMF program page	Where new profiles get posted as they're released	nist.gov/itl/ai-risk-management-framework

Standards

Source	What it is	Link
ISO/IEC 42001:2023	The certifiable AI management system standard (structured like ISO 27001/9001). The standard text is paywalled by ISO; the page below is the authoritative place to confirm the current edition and scope before citing it.	iso.org

Policy and law-review scholarship

Paper	Author(s) / Venue	One-line takeaway	Link
The Brussels Effect	Anu Bradford, Northwestern University Law Review, 2012	The foundational argument for <i>why</i> the EU can set de facto global tech rules through market access alone — the intellectual frame for "the EU AI Act matters even if you never ship there directly"	Northwestern Scholarly Commons
The Brussels Effect (book-length treatment)	Anu Bradford, Oxford University Press, 2020	Full-length version of the above, with AI/tech regulation as a running case study	Oxford Academic
Working paper / SSRN copy of "The Brussels Effect"	Anu Bradford	Freely downloadable working-paper version if the journal link above is unavailable	SSRN

How to read a regulatory text (30 minutes, adapted from the ML-paper method)

1. **Pass 1 (10 min):** recitals (the "whereas" clauses) and definitions article. Ask: what problem is this solving, and what does it call the things I care about?
2. **Pass 2 (15 min):** the operative articles relevant to your question, plus any annex they reference. Ask: who exactly does this bind (provider? deployer? both?), and what's the actual trigger condition?
3. **Pass 3 (5 min):** the transitional/entry-into-force provisions. Ask: is this even in force yet, and has it been amended since the copy I'm reading was published? This pass is the one people skip and the one that gets them in trouble with a fast-moving law like this one.

Keep a one-paragraph note per source with the date you checked it — for this module more than any other in the track, that date is load-bearing information, not metadata.

Resources — AI Compliance

This module ages faster than any other in the track. The most useful thing on this page is the "**how to stay current**" list at the bottom — treat everything above it as a snapshot, not a subscription.

Official trackers and primary sources (bookmark these, not a summary of them)

Source	What it's for	Link
EU AI Act Service Desk	Official European Commission Q&A and implementation guidance	https://ai-act-service-desk.ec.europa.eu/
EU AI Act Implementation Timeline	Independently maintained, closely tracks official dates and amendments	https://artificialintelligenceact.eu/implementation-timeline/
EUR-Lex — Regulation (EU) 2024/1689	The actual binding legal text, all EU languages	https://eur-lex.europa.eu/eli/reg/2024/1689/oj
European Commission — AI Act policy page	The Commission's own framing and links to related instruments	https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai
NIST AI Risk Management Framework	Core framework + profiles (Generative AI, and others as released)	https://www.nist.gov/itl/ai-risk-management-framework
FTC — AI enforcement	Enforcement actions and policy statements, posted directly	https://www.ftc.gov/industry/technology/artificial-intelligence
IAPP US State AI Governance Legislation Tracker	The best single view of the fastest-moving part of the US landscape	https://iapp.org/resources/article/us-state-ai-governance-legislation-tracker
IAPP Global AI Law and Policy Tracker	Same idea, worldwide	https://iapp.org/resources/article/global-ai-legislation-tracker
ISO/IEC 42001:2023	Official standard page — confirm current edition before citing	https://www.iso.org/standard/81230.html
FDA — AI/ML-Enabled Medical Devices	FDA's living page on SaMD AI/ML guidance, including PCCP	https://www.fda.gov/medical-devices/software-medical-device-samd/artificial-intelligence-and-machine-learning-aiml-enabled-medical-devices
Federal Reserve Supervisory Letters (SR)	Where SR 11-7 / SR 26-2 and other bank guidance is published directly	https://www.federalreserve.gov/supervisionreg/srletters/srletters.htm
NYC Department of Consumer and Worker Protection — Local Law 144	Official page for the AEDT bias-audit law, including filed audit summaries	https://www.nyc.gov/site/dca/about/automated-employment-decision-tools.page

Courses

Course	Who it's for	Link
IAPP AI Governance Professional (AIGP) certification prep	The closest thing to an industry-standard credential for this exact job	https://iapp.org/certify/aigp/
Coursera — AI Governance / EU AI Act specializations (search current offerings; providers rotate)	Structured intro for engineers new to the legal side	https://www.coursera.org/
NIST AI RMF playbook workshops / webinars	Free, direct from the framework's authors	https://www.nist.gov/itl/ai-risk-management-framework
Future of Privacy Forum training and reports	Deep, practitioner-oriented on data protection + AI overlap	https://fpf.org/

Communities and ongoing reading

Resource	Why	Link
IAPP (International Association of Privacy Professionals)	The professional home base for this entire discipline — news, trackers, local chapters	https://iapp.org/
r/AIPolicy / r/artificial policy threads	Faster than law-firm blogs for "did anyone else see this filing"	https://www.reddit.com/r/AIPolicy/
Tech Policy Press	Independent journalism specifically on AI/tech regulation, less law-firm-marketing-flavored than most alternatives	https://www.techpolicy.press/
Law firm client alerts (use plural, cross-check)	Fast, free, but incentivized to be alarming — read at least two on any big development before believing either	e.g. Gibson Dunn, Wilson Sonsini, DWT, Hogan Lovells AI practice pages

How to stay current — the habit, not just the bookmarks

This list matters more than any single link above. Before you state a regulatory fact as settled in a design doc or a compliance sign-off:

1. **Check the primary source directly**, not a summary of it — see the "Official trackers" table above. Summaries lag by weeks to months and compound each other's errors once one gets it slightly wrong.
2. **Ask what changed since your last check**, not just "what does it say now" — a law that was comprehensive six months ago and is narrow-and-delayed today (Colorado, notes/01) is a different compliance program, not an update to the old one.
3. **Track litigation, not just legislation** — a statute's text and its enforceable reality can diverge (the EO 14365 preemption fight, the NYT v. OpenAI case) and the gap is exactly where "I read the law" stops being sufficient.
4. **Re-run this module's two tools against a current use case on a quarterly cadence**, the same way you'd re-run a security scan — a classifier or mapper that's never re-exercised silently drifts out of sync with both the law and your own product surface.

5. **When two sources disagree, prefer the one closer to the regulator** — the Service Desk over a law firm blog, the Federal Reserve's own SR-letter page over a summary of it, EUR-Lex over a tracker site (trackers are excellent for *finding* what to check, not for *citing* as the final word).

See also ../12-ai-governance/resources/RESOURCES.md for the internal-governance side of this same reading list, and ../15-ai-evals/ for the evaluation tooling this module's obligations point back to.